

# CMPE 156/L, Winter 2019

## Programming Assignment 4

Due: \_\_\_\_\_

The problem in this assignment is to develop the client and server sides of an application that enables a client to download “chunks” of a file from multiple servers (for example, ftp mirrors) distributed over the Internet, and assemble the chunks to form the complete file. When the server is the bottleneck, this can speed up the download of a large file over normal ftp.

The problem in this assignment is to develop the same parallel file transfer application of the last assignment, using UDP sockets instead of TCP. The application should enable a client to download “chunks” of a file from multiple servers (for example, ftp mirrors) distributed over the Internet, and assemble the chunks to form the complete file. This application has two key differences with the one that used TCP.

1. Each chunk must be transferred from a server to the client using UDP as the transport-layer protocol. Because the chunks must be transferred reliably, this means that you need to implement a reliable application-layer protocol over UDP. The TFTP protocol is a good example of such an application-layer protocol. TFTP uses the stop-and-wait protocol with a retransmission timer for reliability, and block-level sequence numbers to maintain order. Some of the more complex functions of TCP, such as flow control, congestion control and round-trip delay estimation are not required to be supported.
2. The server must support concurrent transfers of chunks of a file to the same client or to different clients. This requires implementing a scheme similar to that employed by TFTP, where the server receives the initial request on a well-known UDP port but opens a separate UDP port for transferring data from/to each client.

To test this application, you will have a single client process on your machine transferring data from multiple servers. The design should allow the server processes to run on any machine on the Internet, although for testing you can run several server processes on the same machine. The client will obtain the contact information for the servers from a text file `server-info.text` that lists the IPv4 addresses of the server hosts and the corresponding server port numbers. The port numbers can be the same if the server processes run on different machines. Otherwise, the only way to distinguish multiple server processes on the same machine is to use different port numbers.

|             |      |
|-------------|------|
| 128.32.16.1 | 4321 |
| 192.71.62.5 | 4321 |
| 127.0.0.1   | 5555 |

The IP address 127.0.0.1 represents “localhost,” indicating that the server is on the local machine. You should start each copy of the server using the command line.

```
./myserver <port-number>
```

where the argument is its listening port number. The client process is started with the command

```
./myclient <server-info.txt> <num-chunks> <filename>
```

The client should first contact one of the servers listed in <server-info.txt> and obtain the size of the file, named <filename>, to be transferred. The filename can contain a directory path. The client should then connect the number of servers as specified in the <num-chunks> argument and transfer a chunk of the file from each server. Each chunk is specified by a starting offset and a chunk size. The client must calculate these parameters by dividing the size of the file by the number of connections. The client must then initiate the connections and transfer each chunk concurrently. When all the chunks have been received, it should assemble them into the complete file and save it on disk (so that you can do diff with the original file).

When the number of chunks <num-chunks> is larger than the number of entries in server-info.txt, the client process needs to transfer more than one chunk from the same server. For example, if the number of chunks is 8, and the number of entries in server-info.txt is only 6, you need to transfer two chunks in parallel from two of the servers, and one chunk each from the remaining four servers.

You should not assume that each chunk will fit within a single UDP segment, so your application layer protocol must deal with breaking a chunk into multiple UDP segments for transmission and re-assembling the UDP segments into the original chunk on the receive side.

## Requirements

- The server must be designed to support concurrent transfers to either the same client or different clients in parallel.
- You need to specify the application protocol, including the commands, responses, message formats, etc., to be used by the client and server to achieve the desired functionality. At a minimum, the client needs to check if a certain file exists at the server and get its size. The client should also be able to initiate the transfer of a chunk of the file, identified by a starting offset and chunk size. You should also take care of any error scenarios that can arise in the client-server interactions.
- If one of the selected servers happens to be unreachable or malfunctions, the client must use an alternate server to transfer the corresponding chunk. If all the servers in server-info.txt are already in use, the client must start the transfer of a new chunk from one of the servers that is currently performing a transfer.
- The server design must be robust against errors and network failures. The application should not fail unless the servers listed in server-info.txt have all failed.
- The server should have a number of files available to serve (at least 6). The files can be limited to text files, and their total sizes can be limited to 100 Kbytes. These files are to be placed in the same directory where the server program is.

## What to submit?

You must submit all the files in a single compressed tar file (with tar.gz extension). The files should include

1. All source code necessary to build and run the client and server.
2. A README file including your name and a list of files in the submission with a brief description of each. If your code does not work completely, explain what works and what doesn't or has not been tested.
3. A Makefile that can be used to build the client and server binaries.
4. Documentation of your design in plain text or pdf. Do not include any Microsoft Word files. The documentation should describe the internal design of your client and server implementations.
5. Including in the design document, a specification of the application-layer protocol, describing the handshakes involved, message formats, error handling, etc.
6. Organize the files into directories (src, bin, doc, etc.)

## Grading

Each submission will be tested to make sure it works properly and can deal with errors. Grades are allocated using the following guidelines:

|                            |     |
|----------------------------|-----|
| Basic functionality        | 20% |
| Basic testing              | 20% |
| Dealing with errors        | 30% |
| Documentation              | 15% |
| Style/Code structure, etc. | 15% |

Note that 30% of the grade will be based on how well your code deals with errors. Good practices include checking all system calls for errors and avoiding unsafe situations such as a buffer overflow.

## Honor Code

The files must be developed independently. All the work submitted must be your own.